

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih Koray Kavukcuoglu David Silver Alex Graves Ioannis Antonoglou

Daan Wierstra Martin Riedmiller

DeepMind Technologies

{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com

Abstract

We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The model is a convolutional neural network, trained with a variant of Q-learning, whose input is raw pixels and whose output is a value function estimating future rewards. We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them.

1 Introduction

Learning to control agents directly from high-dimensional sensory inputs like vision and speech is one of the long-standing challenges of reinforcement learning (RL). Most successful RL applications that operate on these domains have relied on hand-crafted features combined with linear value functions or policy representations. Clearly, the performance of such systems heavily relies on the quality of the feature representation.

Recent advances in deep learning have made it possible to extract high-level features from raw sensory data, leading to breakthroughs in computer vision [11, 22, 16] and speech recognition [6, 7]. These methods utilise a range of neural network architectures, including convolutional networks, multilayer perceptrons, restricted Boltzmann machines and recurrent neural networks, and have exploited both supervised and unsupervised learning. It seems natural to ask whether similar techniques could also be beneficial for RL with sensory data.

However reinforcement learning presents several challenges from a deep learning perspective. Firstly, most successful deep learning applications to date have required large amounts of hand-labelled training data. RL algorithms, on the other hand, must be able to learn from a scalar reward signal that is frequently sparse, noisy and delayed. The delay between actions and resulting rewards, which can be thousands of timesteps long, seems particularly daunting when compared to the direct association between inputs and targets found in supervised learning. Another issue is that most deep learning algorithms assume the data samples to be independent, while in reinforcement learning one typically encounters sequences of highly correlated states. Furthermore, in RL the data distribution changes as the algorithm learns new behaviours, which can be problematic for deep learning methods that assume a fixed underlying distribution.

This paper demonstrates that a convolutional neural network can overcome these challenges to learn successful control policies from raw video data in complex RL environments. The network is trained with a variant of the Q-learning [26] algorithm, with stochastic gradient descent to update the weights. To alleviate the problems of correlated data and non-stationary distributions, we use

使用深度强化学习玩Atari游戏

弗拉基米尔·米尼 科拉伊·卡武克屈奥卢 大卫·西尔弗 亚历克斯·格雷夫斯 约安尼斯·安东诺格鲁

达安·维尔斯特拉 马丁·里德米勒

深度思维科技

{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com

摘要

我们提出了首个通过强化学习直接从高维度感官输入中成功学习控制策略的深度学习模型。该模型是一个卷积神经网络，采用Q-learning的变体进行训练，其输入为原始像素，输出为用于估计未来奖励的价值函数。我们将该方法应用于街机学习环境中的七款Atari 2600游戏，未对架构或学习算法进行任何调整。研究发现，该方法在六款游戏中超越了所有先前方案，并在其中三款游戏中超越了人类专家水平。

1 引言

学习如何直接从高维感官输入（如视觉和语音）中控制智能体，是强化学习（RL）长期面临的挑战之一。在这些领域运行的大多数成功RL应用，都依赖于手工设计的特征与线性价值函数或策略表示相结合。显然，这类系统的性能在很大程度上取决于特征表示的质量。

深度学习的最新进展使得从原始感官数据中提取高级特征成为可能，从而在计算机视觉[11, 22, 16]和语音识别[6, 7]领域取得了突破。这些方法利用了包括卷积网络、多层感知器、受限玻尔兹曼机和循环神经网络在内的一系列神经网络架构，并同时运用了监督学习和无监督学习。很自然地，我们会问类似的技术是否也能对基于感官数据的强化学习（RL）带来益处。

然而，从深度学习的角度来看，强化学习提出了若干挑战。首先，迄今为止大多数成功的深度学习应用都需要大量手工标注的训练数据。而强化学习算法必须能够从通常稀疏、嘈杂且延迟的标量奖励信号中学习。行动与随之而来的奖励之间的延迟可能长达数千个时间步，与监督学习中输入与目标之间的直接关联相比，这显得尤为艰巨。另一个问题是，大多数深度学习算法假设数据样本是独立的，而在强化学习中，通常会遇到高度相关的状态序列。此外，在强化学习中，数据分布会随着算法学习新行为而改变，这对于假设底层分布固定的深度学习方法来说可能是个问题。

本文证明，卷积神经网络能够克服这些挑战，在复杂的强化学习环境中从原始视频数据中学习成功的控制策略。该网络采用Q学习[26]算法的变体进行训练，并通过随机梯度下降更新权重。为缓解数据相关性和非平稳分布问题，我们采用




Figure 1: Screen shots from five Atari 2600 Games: (Left-to-right) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

an experience replay mechanism [13] which randomly samples previous transitions, and thereby smooths the training distribution over many past behaviors.

We apply our approach to a range of Atari 2600 games implemented in The Arcade Learning Environment (ALE) [3]. Atari 2600 is a challenging RL testbed that presents agents with a high dimensional visual input (210×160 RGB video at 60Hz) and a diverse and interesting set of tasks that were designed to be difficult for humans players. Our goal is to create a single neural network agent that is able to successfully learn to play as many of the games as possible. The network was not provided with any game-specific information or hand-designed visual features, and was not privy to the internal state of the emulator; it learned from nothing but the video input, the reward and terminal signals, and the set of possible actions—just as a human player would. Furthermore the network architecture and all hyperparameters used for training were kept constant across the games. So far the network has outperformed all previous RL algorithms on six of the seven games we have attempted and surpassed an expert human player on three of them. Figure 1 provides sample screenshots from five of the games used for training.

2 Background

We consider tasks in which an agent interacts with an environment $\mathcal{E}$, in this case the Atari emulator, in a sequence of actions, observations and rewards. At each time-step the agent selects an action a_t from the set of legal game actions, $\mathcal{A} = \{1, \dots, K\}$. The action is passed to the emulator and modifies its internal state and the game score. In general $\mathcal{E}$ may be stochastic. The emulator’s internal state is not observed by the agent; instead it observes an image $x_t \in \mathbb{R}^d$ from the emulator, which is a vector of raw pixel values representing the current screen. In addition it receives a reward r_t representing the change in game score. Note that in general the game score may depend on the whole prior sequence of actions and observations; feedback about an action may only be received after many thousands of time-steps have elapsed.

Since the agent only observes images of the current screen, the task is partially observed and many emulator states are perceptually aliased, i.e. it is impossible to fully understand the current situation from only the current screen x_t . We therefore consider sequences of actions and observations, $s_t = x_1, a_1, x_2, \dots, a_{t-1}, x_t$, and learn game strategies that depend upon these sequences. All sequences in the emulator are assumed to terminate in a finite number of time-steps. This formalism gives rise to a large but finite Markov decision process (MDP) in which each sequence is a distinct state. As a result, we can apply standard reinforcement learning methods for MDPs, simply by using the complete sequence s_t as the state representation at time t .

The goal of the agent is to interact with the emulator by selecting actions in a way that maximises future rewards. We make the standard assumption that future rewards are discounted by a factor of γ per time-step, and define the future discounted *return* at time t as $R_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$, where T is the time-step at which the game terminates. We define the optimal action-value function $Q^*(s, a)$ as the maximum expected return achievable by following any strategy, after seeing some sequence s and then taking some action a , $Q^*(s, a) = \max_{\pi} \mathbb{E}[R_t | s_t = s, a_t = a, \pi]$, where π is a policy mapping sequences to actions (or distributions over actions).

The optimal action-value function obeys an important identity known as the *Bellman equation*. This is based on the following intuition: if the optimal value $Q^*(s', a')$ of the sequence s' at the next time-step was known for all possible actions a' , then the optimal strategy is to select the action a'




图1: 五款Atari 2600游戏的屏幕截图: (Left-to-right) 《乒乓》 《打砖块》 《太空侵略者》 《海底冒险》 《光束骑士》

一种经验回放机制[13], 它随机采样先前的转换, 从而平滑了基于许多过去行为的训练分布。

我们将我们的方法应用于一系列在街机学习环境 (ALE) [3]中实现的Atari 2600游戏。Atari 2600是一个具有挑战性的强化学习测试平台, 它为智能体提供了高维度的视觉输入 (210×160 RGB视频, 每秒60帧) 以及一系列多样且有趣的任务, 这些任务原本设计得对人类玩家来说颇具难度。我们的目标是创建一个单一的神经网络智能体, 能够成功学习并尽可能多地掌握这些游戏。该网络未获得任何游戏特定信息或手动设计的视觉特征, 也无法访问模拟器的内部状态; 它仅从视频输入、奖励与终止信号以及可能的动作集合中学习——正如人类玩家一样。此外, 网络架构和所有用于训练的超参数在不同游戏中均保持不变。迄今为止, 在我们尝试的七款游戏中, 该网络已在六款上超越了所有先前的强化学习算法, 并在其中三款游戏中超越了人类专家玩家。图1展示了用于训练的五款游戏的示例截图。

2 背景

我们考虑的任务中, 智能体与环境 $\mathcal{E}$ (此处指Atari模拟器) 通过一系列动作、观察和奖励进行交互。在每个时间步, 智能体从合法游戏动作集合 $\mathcal{A} = \{1, \dots, K\}$ 中选择一个动作 a_t 。该动作被传递至模拟器, 并改变其内部状态和游戏分数。通常 $\mathcal{E}$ 可能是随机的。智能体无法观测模拟器的内部状态, 而是观测到来自模拟器的一幅图像 $x_t \in \mathbb{R}^d$, 这是一个代表当前屏幕的原始像素值向量。此外, 智能体会接收到一个奖励 r_t , 代表游戏分数的变化。需要注意的是, 游戏分数通常可能依赖于整个先前的动作和观察序列; 关于某个动作的反馈可能要在数千个时间步之后才能收到。

由于智能体仅能观察到当前屏幕的图像, 该任务属于部分可观测, 且许多模拟器状态在感知上存在歧义, 即仅凭当前屏幕 x_t 无法完全理解当前状况。因此, 我们考虑动作与观测的序列 $s_t = x_1, a_1, x_2, \dots, a_{t-1}, x_t$, 并学习依赖于这些序列的游戏策略。假设模拟器中的所有序列都在有限时间步内终止。这种形式化方法产生了一个庞大但有限的马尔可夫决策过程 (MDP), 其中每个序列都是一个独立的状态。因此, 我们可以通过将完整序列 s_t 作为时间 t 的状态表示, 直接应用标准的MDP强化学习方法。

智能体的目标是通过选择行动与模拟器互动, 以最大化未来奖励。我们采用标准假设, 即未来奖励在每个时间步按系数 γ 进行折现, 并将时间 t 的未来折现*return*定义为 $R_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$, 其中 T 是游戏终止的时间步。我们将最优行动价值函数 $Q^*(s, a)$ 定义为在看到某个序列 s 后采取某个行动 a , 并遵循任意策略所能获得的最大期望回报, 即 $Q^*(s, a) = \max_{\pi} \mathbb{E}[R_t | s_t = s, a_t = a, \pi]$, 其中 π 是将序列映射到行动 (或行动分布) 的策略。

最优动作价值函数遵循一个重要的恒等式, 称为*Bellman equation*。其基本思想是: 如果已知下一时间步所有可能动作 a' 对应的序列 s' 的最优值 $Q^*(s', a')$, 那么最优策略就是选择动作 a'

maximising the expected value of $r + \gamma Q^*(s', a')$,

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \middle| s, a \right] \quad (1)$$

The basic idea behind many reinforcement learning algorithms is to estimate the action-value function, by using the Bellman equation as an iterative update, $Q_{i+1}(s, a) = \mathbb{E} [r + \gamma \max_{a'} Q_i(s', a') | s, a]$. Such *value iteration* algorithms converge to the optimal action-value function, $Q_i \rightarrow Q^*$ as $i \rightarrow \infty$ [23]. In practice, this basic approach is totally impractical, because the action-value function is estimated separately for each sequence, without any generalisation. Instead, it is common to use a function approximator to estimate the action-value function, $Q(s, a; \theta) \approx Q^*(s, a)$. In the reinforcement learning community this is typically a linear function approximator, but sometimes a non-linear function approximator is used instead, such as a neural network. We refer to a neural network function approximator with weights θ as a Q-network. A Q-network can be trained by minimising a sequence of loss functions $L_i(\theta_i)$ that changes at each iteration i ,

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right], \quad (2)$$

where $y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$ is the target for iteration i and $\rho(s, a)$ is a probability distribution over sequences s and actions a that we refer to as the *behaviour distribution*. The parameters from the previous iteration θ_{i-1} are held fixed when optimising the loss function $L_i(\theta_i)$. Note that the targets depend on the network weights; this is in contrast with the targets used for supervised learning, which are fixed before learning begins. Differentiating the loss function with respect to the weights we arrive at the following gradient,

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]. \quad (3)$$

Rather than computing the full expectations in the above gradient, it is often computationally expedient to optimise the loss function by stochastic gradient descent. If the weights are updated after every time-step, and the expectations are replaced by single samples from the behaviour distribution ρ and the emulator $\mathcal{E}$ respectively, then we arrive at the familiar *Q-learning* algorithm [26].

Note that this algorithm is *model-free*: it solves the reinforcement learning task directly using samples from the emulator $\mathcal{E}$, without explicitly constructing an estimate of $\mathcal{E}$. It is also *off-policy*: it learns about the greedy strategy $a = \max_a Q(s, a; \theta)$, while following a behaviour distribution that ensures adequate exploration of the state space. In practice, the behaviour distribution is often selected by an ϵ -greedy strategy that follows the greedy strategy with probability $1 - \epsilon$ and selects a random action with probability ϵ .

3 Related Work

Perhaps the best-known success story of reinforcement learning is *TD-gammon*, a backgammon-playing program which learnt entirely by reinforcement learning and self-play, and achieved a super-human level of play [24]. TD-gammon used a model-free reinforcement learning algorithm similar to Q-learning, and approximated the value function using a multi-layer perceptron with one hidden layer¹.

However, early attempts to follow up on TD-gammon, including applications of the same method to chess, Go and checkers were less successful. This led to a widespread belief that the TD-gammon approach was a special case that only worked in backgammon, perhaps because the stochasticity in the dice rolls helps explore the state space and also makes the value function particularly smooth [19].

Furthermore, it was shown that combining model-free reinforcement learning algorithms such as Q-learning with non-linear function approximators [25], or indeed with off-policy learning [1] could cause the Q-network to diverge. Subsequently, the majority of work in reinforcement learning focused on linear function approximators with better convergence guarantees [25].

¹In fact TD-Gammon approximated the state value function $V(s)$ rather than the action-value function $Q(s, a)$, and learnt *on-policy* directly from the self-play games

最大化 $r + \gamma Q^*(s', a')$ 的期望值,

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \middle| s, a \right] \quad (1)$$

许多强化学习算法的基本思想是通过使用贝尔曼方程作为迭代更新来估计动作价值函数, $Q_{i+1}(s, a) = \mathbb{E} [r + \gamma \max_{a'} Q_i(s', a') | s, a]$ 。这样的 *value iteration* 算法会收敛到最优动作价值函数, $Q_i \rightarrow Q^*$ 如 $i \rightarrow \infty$ [23]所述。在实践中, 这种基本方法完全不切实际, 因为动作价值函数是针对每个序列单独估计的, 没有任何泛化能力。相反, 通常使用函数逼近器来估计动作价值函数, $Q(s, a; \theta) \approx Q^*(s, a)$ 。在强化学习领域, 这通常是一个线性函数逼近器, 但有时也会使用非线性函数逼近器, 例如神经网络。我们将具有权重 θ 的神经网络函数逼近器称为Q网络。可以通过最小化一系列在每次迭代 i 时变化的损失函数 $L_i(\theta_i)$ 来训练Q网络。

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right], \quad (2)$$

其中 $y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$ 是迭代 i 的目标, $\rho(s, a)$ 是序列 s 和动作 a 上的概率分布, 我们称之为 *behaviour distribution*。在优化损失函数 $L_i(\theta_i)$ 时, 前一次迭代 θ_{i-1} 的参数保持不变。注意, 目标值依赖于网络权重; 这与监督学习使用的目标不同, 后者在学习开始前就已固定。对损失函数关于权重求导, 我们得到以下梯度,

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]. \quad (3)$$

与其计算上述梯度中的完整期望, 通常采用随机梯度下降优化损失函数在计算上更为便利。如果在每个时间步后更新权重, 并将期望分别替换为行为分布 ρ 和模拟器 $\mathcal{E}$ 的单个样本, 那么我们就得到了熟悉的 *Q-learning* 算法[26]。

请注意, 该算法是 *model-free*: 它直接使用来自模拟器 $\mathcal{E}$ 的样本来解决强化学习任务, 而无需显式构建 $\mathcal{E}$ 的估计。它也是 *off-policy*: 在学习贪婪策略 $a = \max_a Q(s, a; \theta)$ 的同时, 遵循一个确保充分探索状态空间的行为分布。在实践中, 行为分布通常通过 ϵ -贪婪策略来选择, 该策略以概率 $1 - \epsilon$ 遵循贪婪策略, 并以概率 ϵ 选择随机动作。

3 相关工作

强化学习最著名的成功案例或许是 *TD-gammon*, 这是一个完全通过强化学习和自我对弈学习的西洋双陆棋程序, 达到了超人类的游戏水平[24]。TD-gammon采用了类似Q学习的无模型强化学习算法, 并使用带有一个隐藏层的多层感知器来逼近价值函数¹。

然而, 早期对TD-gammon的后续尝试, 包括将相同方法应用于国际象棋、围棋和跳棋, 都未能取得显著成功。这导致了一种普遍看法, 认为TD-gammon方法只是一个特例, 仅适用于双陆棋——或许是因为骰子投掷的随机性有助于探索状态空间, 同时也使价值函数特别平滑[19]。

此外, 研究表明, 将无模型强化学习算法(如Q学习)与非线性函数近似器[25]结合, 或者实际上与离策略学习[1]结合, 可能导致Q网络发散。随后, 强化学习领域的大部分工作都集中在具有更好收敛保证的线性函数近似器上[25]。

¹In fact TD-Gammon approximated the state value function $V(s)$ rather than the action-value function $Q(s, a)$, and learnt *on-policy* directly from the self-play games

More recently, there has been a revival of interest in combining deep learning with reinforcement learning. Deep neural networks have been used to estimate the environment $\mathcal{E}$; restricted Boltzmann machines have been used to estimate the value function [21]; or the policy [9]. In addition, the divergence issues with Q-learning have been partially addressed by *gradient temporal-difference* methods. These methods are proven to converge when evaluating a fixed policy with a nonlinear function approximator [14]; or when learning a control policy with linear function approximation using a restricted variant of Q-learning [15]. However, these methods have not yet been extended to nonlinear control.

Perhaps the most similar prior work to our own approach is neural fitted Q-learning (NFQ) [20]. NFQ optimises the sequence of loss functions in Equation 2, using the RPROP algorithm to update the parameters of the Q-network. However, it uses a batch update that has a computational cost per iteration that is proportional to the size of the data set, whereas we consider stochastic gradient updates that have a low constant cost per iteration and scale to large data-sets. NFQ has also been successfully applied to simple real-world control tasks using purely visual input, by first using deep autoencoders to learn a low dimensional representation of the task, and then applying NFQ to this representation [12]. In contrast our approach applies reinforcement learning end-to-end, directly from the visual inputs; as a result it may learn features that are directly relevant to discriminating action-values. Q-learning has also previously been combined with experience replay and a simple neural network [13], but again starting with a low-dimensional state rather than raw visual inputs.

The use of the Atari 2600 emulator as a reinforcement learning platform was introduced by [3], who applied standard reinforcement learning algorithms with linear function approximation and generic visual features. Subsequently, results were improved by using a larger number of features, and using tug-of-war hashing to randomly project the features into a lower-dimensional space [2]. The HyperNEAT evolutionary architecture [8] has also been applied to the Atari platform, where it was used to evolve (separately, for each distinct game) a neural network representing a strategy for that game. When trained repeatedly against deterministic sequences using the emulator’s reset facility, these strategies were able to exploit design flaws in several Atari games.

4 Deep Reinforcement Learning

Recent breakthroughs in computer vision and speech recognition have relied on efficiently training deep neural networks on very large training sets. The most successful approaches are trained directly from the raw inputs, using lightweight updates based on stochastic gradient descent. By feeding sufficient data into deep neural networks, it is often possible to learn better representations than handcrafted features [11]. These successes motivate our approach to reinforcement learning. Our goal is to connect a reinforcement learning algorithm to a deep neural network which operates directly on RGB images and efficiently process training data by using stochastic gradient updates.

Tesauro’s TD-Gammon architecture provides a starting point for such an approach. This architecture updates the parameters of a network that estimates the value function, directly from on-policy samples of experience, $s_t, a_t, r_t, s_{t+1}, a_{t+1}$, drawn from the algorithm’s interactions with the environment (or by self-play, in the case of backgammon). Since this approach was able to outperform the best human backgammon players 20 years ago, it is natural to wonder whether two decades of hardware improvements, coupled with modern deep neural network architectures and scalable RL algorithms might produce significant progress.

In contrast to TD-Gammon and similar online approaches, we utilize a technique known as *experience replay* [13] where we store the agent’s experiences at each time-step, $e_t = (s_t, a_t, r_t, s_{t+1})$ in a data-set $\mathcal{D} = e_1, \dots, e_N$, pooled over many episodes into a *replay memory*. During the inner loop of the algorithm, we apply Q-learning updates, or minibatch updates, to samples of experience, $e \sim \mathcal{D}$, drawn at random from the pool of stored samples. After performing experience replay, the agent selects and executes an action according to an ϵ -greedy policy. Since using histories of arbitrary length as inputs to a neural network can be difficult, our Q-function instead works on fixed length representation of histories produced by a function ϕ . The full algorithm, which we call *deep Q-learning*, is presented in Algorithm 1.

This approach has several advantages over standard online Q-learning [23]. First, each step of experience is potentially used in many weight updates, which allows for greater data efficiency.

最近，将深度学习与强化学习相结合的兴趣再度兴起。深度神经网络已被用于估计环境 $\mathcal{E}$ ；受限玻尔兹曼机被用于估计价值函数[21]或策略[9]。此外，Q学习的发散问题已通过 *gradient temporal-difference*方法得到部分解决。这些方法在评估固定策略时使用非线性函数逼近器[14]，或在使用Q学习受限变体进行线性函数逼近学习控制策略时[15]，已被证明能够收敛。然而，这些方法尚未扩展到非线性控制领域。

或许与我们方法最为相似的先前工作是神经拟合Q学习（NFQ）[20]。NFQ优化了公式2中的损失函数序列，使用RPROP算法更新Q网络的参数。然而，它采用批量更新方式，每次迭代的计算成本与数据集大小成正比；而我们考虑使用随机梯度更新，每次迭代成本较低且能扩展至大型数据集。NFQ还通过先利用深度自编码器学习任务的低维表示，再对该表示应用NFQ，成功应用于仅依赖视觉输入的简单现实控制任务[12]。相比之下，我们的方法从视觉输入端到端地直接应用强化学习；因此可能学习到与区分动作价值直接相关的特征。Q学习此前也曾与经验回放及简单神经网络结合[13]，但同样是从低维状态而非原始视觉输入开始。

[3]首次将Atari 2600模拟器作为强化学习平台使用，他们采用线性函数逼近和通用视觉特征的标准强化学习算法。随后，通过使用更多特征，并利用河哈希将特征随机投影到低维空间[2]，结果得到了改善。HyperNEAT进化架构[8]也被应用于Atari平台，用于（针对每个不同的游戏分别）进化出代表该游戏策略的神经网络。当通过模拟器的重置功能反复针对确定性序列进行训练时，这些策略能够利用多个Atari游戏中的设计缺陷。

4 深度强化学习

计算机视觉和语音识别领域的最新突破依赖于在非常大的训练集上高效训练深度神经网络。最成功的方法直接从原始输入进行训练，使用基于随机梯度下降的轻量级更新。通过向深度神经网络输入足够的数据，通常可以学习到比手工特征更好的表示[11]。这些成功激励了我们在强化学习方面的探索。我们的目标是将强化学习算法与一个深度神经网络连接起来，该网络直接处理RGB图像，并通过使用随机梯度更新高效处理训练数据。

Tesauro的TD-Gammon架构为这种方法提供了一个起点。该架构通过直接从算法与环境交互（或在双陆棋案例中通过自我对弈）产生的策略样本 $s_t, a_t, r_t, s_{t+1}, a_{t+1}$ 中，更新估计价值函数的网络参数。由于这种方法在20年前就已能超越最优秀的人类双陆棋选手，人们自然会思考：二十年的硬件进步，结合现代深度神经网络架构与可扩展的强化学习算法，是否能带来重大突破。

与TD-Gammon及类似的在线方法不同，我们采用了一种称为*experience replay*[13]的技术，该技术将智能体在每个时间步的经验存储于数据集 $\mathcal{D} = e_1, \dots, e_N$ 中，并通过多轮次汇总至 *replay memory*。在算法的内部循环中，我们对从存储样本池中随机抽取的经验样本 $e \sim \mathcal{D}$ 应用Q学习更新或小批量更新。完成经验回放后，智能体根据 ϵ -贪婪策略选择并执行动作。由于将任意长度的历史记录作为神经网络输入可能较为困难，我们的Q函数转而处理由函数 ϕ 生成的固定长度历史表示。我们将完整算法称为*deep Q-learning*，其具体流程如算法1所示。

这种方法相较于标准的在线Q学习[23]具有多项优势。首先，每一步经验都可能用于多次权重更新，从而实现了更高的数据效率。

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory $\mathcal{D}$ to capacity N
Initialize action-value function Q with random weights
for episode = 1, M **do**
 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$
 for $t = 1, T$ **do**
 With probability ϵ select a random action a_t
 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$
 Execute action a_t in emulator and observe reward r_t and image x_{t+1}
 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$
 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in $\mathcal{D}$
 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from $\mathcal{D}$
 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3
 end for
end for

Second, learning directly from consecutive samples is inefficient, due to the strong correlations between the samples; randomizing the samples breaks these correlations and therefore reduces the variance of the updates. Third, when learning on-policy the current parameters determine the next data sample that the parameters are trained on. For example, if the maximizing action is to move left then the training samples will be dominated by samples from the left-hand side; if the maximizing action then switches to the right then the training distribution will also switch. It is easy to see how unwanted feedback loops may arise and the parameters could get stuck in a poor local minimum, or even diverge catastrophically [25]. By using experience replay the behavior distribution is averaged over many of its previous states, smoothing out learning and avoiding oscillations or divergence in the parameters. Note that when learning by experience replay, it is necessary to learn off-policy (because our current parameters are different to those used to generate the sample), which motivates the choice of Q-learning.

In practice, our algorithm only stores the last N experience tuples in the replay memory, and samples uniformly at random from $\mathcal{D}$ when performing updates. This approach is in some respects limited since the memory buffer does not differentiate important transitions and always overwrites with recent transitions due to the finite memory size N . Similarly, the uniform sampling gives equal importance to all transitions in the replay memory. A more sophisticated sampling strategy might emphasize transitions from which we can learn the most, similar to prioritized sweeping [17].

4.1 Preprocessing and Model Architecture

Working directly with raw Atari frames, which are 210×160 pixel images with a 128 color palette, can be computationally demanding, so we apply a basic preprocessing step aimed at reducing the input dimensionality. The raw frames are preprocessed by first converting their RGB representation to gray-scale and down-sampling it to a 110×84 image. The final input representation is obtained by cropping an 84×84 region of the image that roughly captures the playing area. The final cropping stage is only required because we use the GPU implementation of 2D convolutions from [11], which expects square inputs. For the experiments in this paper, the function ϕ from algorithm 1 applies this preprocessing to the last 4 frames of a history and stacks them to produce the input to the Q -function.

There are several possible ways of parameterizing Q using a neural network. Since Q maps history-action pairs to scalar estimates of their Q-value, the history and the action have been used as inputs to the neural network by some previous approaches [20, 12]. The main drawback of this type of architecture is that a separate forward pass is required to compute the Q-value of each action, resulting in a cost that scales linearly with the number of actions. We instead use an architecture in which there is a separate output unit for each possible action, and only the state representation is an input to the neural network. The outputs correspond to the predicted Q-values of the individual action for the input state. The main advantage of this type of architecture is the ability to compute Q-values for all possible actions in a given state with only a single forward pass through the network.

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory $\mathcal{D}$ to capacity N
Initialize action-value function Q with random weights
for episode = 1, M **do**
 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$
 for $t = 1, T$ **do**
 With probability ϵ select a random action a_t
 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$
 Execute action a_t in emulator and observe reward r_t and image x_{t+1}
 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$
 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in $\mathcal{D}$
 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from $\mathcal{D}$
 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3
 end for
end for

其次，直接从连续样本中学习效率低下，因为样本之间存在强相关性；随机化样本打破了这些相关性，从而降低了更新的方差。第三，在策略学习时，当前参数决定了参数训练的下一个数据样本。例如，如果最大化动作是向左移动，那么训练样本将主要由左侧的样本主导；如果最大化动作随后切换到右侧，训练分布也会随之切换。很容易看出，不良的反馈循环可能出现，参数可能陷入较差的局部最小值，甚至灾难性地发散[25]。通过使用经验回放，行为分布在许多先前状态上取平均，从而平滑学习过程，避免参数振荡或发散。需要注意的是，通过经验回放学习时，必须采用离策略学习（因为我们当前的参数与生成样本时使用的参数不同），这促使了Q学习的选择。

在实践中，我们的算法仅将最近的 N 条经验元组存储在回放记忆中，并在执行更新时从 $\mathcal{D}$ 中均匀随机采样。这种方法在某些方面存在局限，因为记忆缓冲区不会区分重要的转换，并且由于有限的内存大小 N ，总是用最近的转换覆盖。同样，均匀采样对回放记忆中的所有转换赋予同等重要性。更复杂的采样策略可能会强调那些我们能从中学习最多的转换，类似于优先扫描[17]。

4.1 预处理与模型架构

直接处理原始的Atari帧（即 210×160 像素、128色调色板的图像）在计算上可能要求较高，因此我们采用了一个基本的预处理步骤，旨在降低输入维度。原始帧的预处理首先将其RGB表示转换为灰度图，并下采样至 110×84 图像。最终输入表示是通过裁剪图像中大致覆盖游戏区域的 84×84 区域获得的。最终裁剪阶段仅因我们使用了[11]中的GPU二维卷积实现而必需，该实现要求输入为方形。对于本文中的实验，算法1中的函数 ϕ 将此预处理应用于历史记录的最后4帧，并将它们堆叠以生成 Q 函数的输入。

有多种可能的方法可以使用神经网络对 Q 进行参数化。由于 Q 将历史-动作对映射到其Q值的标量估计，一些先前的方法[20, 12]已将历史和动作作用神经网络的输入。这类架构的主要缺点是需要单独的前向传递来计算每个动作的Q值，导致成本随动作数量线性增长。我们采用了一种不同的架构，其中每个可能的动作都有一个单独的输出单元，只有状态表示作为神经网络的输入。输出对应于输入状态下各个动作的预测Q值。这类架构的主要优点是能够仅通过一次网络前向传递，就计算出给定状态下所有可能动作的Q值。

We now describe the exact architecture used for all seven Atari games. The input to the neural network consists is an $84 \times 84 \times 4$ image produced by ϕ . The first hidden layer convolves $16 \ 8 \times 8$ filters with stride 4 with the input image and applies a rectifier nonlinearity [10, 18]. The second hidden layer convolves $32 \ 4 \times 4$ filters with stride 2, again followed by a rectifier nonlinearity. The final hidden layer is fully-connected and consists of 256 rectifier units. The output layer is a fully-connected linear layer with a single output for each valid action. The number of valid actions varied between 4 and 18 on the games we considered. We refer to convolutional networks trained with our approach as Deep Q-Networks (DQN).

5 Experiments

So far, we have performed experiments on seven popular ATARI games – Beam Rider, Breakout, Enduro, Pong, Q*bert, Seaquest, Space Invaders. We use the same network architecture, learning algorithm and hyperparameters settings across all seven games, showing that our approach is robust enough to work on a variety of games without incorporating game-specific information. While we evaluated our agents on the real and unmodified games, we made one change to the reward structure of the games during training only. Since the scale of scores varies greatly from game to game, we fixed all positive rewards to be 1 and all negative rewards to be -1 , leaving 0 rewards unchanged. Clipping the rewards in this manner limits the scale of the error derivatives and makes it easier to use the same learning rate across multiple games. At the same time, it could affect the performance of our agent since it cannot differentiate between rewards of different magnitude.

In these experiments, we used the RMSProp algorithm with minibatches of size 32. The behavior policy during training was ϵ -greedy with ϵ annealed linearly from 1 to 0.1 over the first million frames, and fixed at 0.1 thereafter. We trained for a total of 10 million frames and used a replay memory of one million most recent frames.

Following previous approaches to playing Atari games, we also use a simple frame-skipping technique [3]. More precisely, the agent sees and selects actions on every k^{th} frame instead of every frame, and its last action is repeated on skipped frames. Since running the emulator forward for one step requires much less computation than having the agent select an action, this technique allows the agent to play roughly k times more games without significantly increasing the runtime. We use $k = 4$ for all games except Space Invaders where we noticed that using $k = 4$ makes the lasers invisible because of the period at which they blink. We used $k = 3$ to make the lasers visible and this change was the only difference in hyperparameter values between any of the games.

5.1 Training and Stability

In supervised learning, one can easily track the performance of a model during training by evaluating it on the training and validation sets. In reinforcement learning, however, accurately evaluating the progress of an agent during training can be challenging. Since our evaluation metric, as suggested by [3], is the total reward the agent collects in an episode or game averaged over a number of games, we periodically compute it during training. The average total reward metric tends to be very noisy because small changes to the weights of a policy can lead to large changes in the distribution of states the policy visits. The leftmost two plots in figure 2 show how the average total reward evolves during training on the games Seaquest and Breakout. Both averaged reward plots are indeed quite noisy, giving one the impression that the learning algorithm is not making steady progress. Another, more stable, metric is the policy’s estimated action-value function Q , which provides an estimate of how much discounted reward the agent can obtain by following its policy from any given state. We collect a fixed set of states by running a random policy before training starts and track the average of the maximum² predicted Q for these states. The two rightmost plots in figure 2 show that average predicted Q increases much more smoothly than the average total reward obtained by the agent and plotting the same metrics on the other five games produces similarly smooth curves. In addition to seeing relatively smooth improvement to predicted Q during training we did not experience any divergence issues in any of our experiments. This suggests that, despite lacking any theoretical convergence guarantees, our method is able to train large neural networks using a reinforcement learning signal and stochastic gradient descent in a stable manner.

²The maximum for each state is taken over the possible actions.

我们现在描述用于所有七款Atari游戏的精确架构。神经网络的输入是由 ϕ 生成的 $84 \times 84 \times 4$ 图像。第一个隐藏层使用16个 8×8 滤波器以步长4对输入图像进行卷积，并应用整流非线性激活函数[10, 18]。第二个隐藏层使用32个 4×4 滤波器以步长2进行卷积，同样后接整流非线性激活。最后的隐藏层是全连接的，包含256个整流单元。输出层是一个全连接的线性层，每个有效动作对应一个输出。在我们考虑的游戏里，有效动作的数量在4到18之间变化。我们将用我们的方法训练的卷积网络称为深度Q网络（DQN）。

5 实验

迄今为止，我们已在七款流行的ATARI游戏上进行了实验——包括《Beam Rider》、《Breakout》、《Enduro》、《Pong》、《Q*bert》、《Seaquest》和《Space Invaders》。我们在所有七款游戏中使用了相同的网络架构、学习算法和超参数设置，这表明我们的方法足够稳健，无需融入游戏特定信息即可适用于多种游戏。虽然我们在真实且未经修改的游戏上评估了智能体，但在训练期间对游戏的奖励结构进行了一项调整。由于不同游戏的得分范围差异巨大，我们将所有正向奖励固定为1，所有负向奖励设为 -1 ，而0奖励保持不变。通过这种方式裁剪奖励，限制了误差导数的范围，使得在多个游戏中使用相同学习率更为容易。然而，这也可能影响智能体的表现，因为它无法区分不同量级的奖励。

在这些实验中，我们使用了RMSProp算法，并采用大小为32的小批量。训练期间的行为策略为 ϵ -贪婪策略，其中 ϵ 在前一百万帧中从1线性退火至0.1，之后固定为0.1。我们总共训练了一千万帧，并使用了一百万帧最近帧的回放记忆。

遵循先前玩Atari游戏的方法，我们也采用了一种简单的跳帧技术[3]。更具体地说，智能体每隔 k^{th} 帧观察并选择动作，而不是每一帧，并且在跳过的帧上重复其上一个动作。由于向前运行模拟器一步所需的计算量远小于让智能体选择动作，这项技术使得智能体能够在不显著增加运行时间的情况下，玩大约 k 倍多的游戏。除了《太空侵略者》外，我们对所有游戏都使用 $k = 4$ ，因为我们注意到在《太空侵略者》中使用 $k = 4$ 会使激光不可见，这是由于它们的闪烁周期。我们使用 $k = 3$ 来使激光可见，这一变化是各游戏间超参数值的唯一差异。

5.1 训练与稳定性

在监督学习中，通过评估模型在训练集和验证集上的表现，可以轻松跟踪训练期间模型的性能。然而，在强化学习中，准确评估智能体在训练过程中的进展可能具有挑战性。由于我们的评估指标（如[3]所建议）是智能体在多轮游戏或单局游戏中获得的总奖励的平均值，我们会在训练期间定期计算该指标。平均总奖励指标往往波动很大，因为策略权重的微小变化可能导致策略访问的状态分布发生巨大改变。图2中最左侧的两个图表展示了《海底探险》和《打砖块》游戏中平均总奖励在训练期间的变化情况。这两个平均奖励图表确实相当波动，给人一种学习算法未取得稳定进展的印象。另一个更稳定的指标是策略的估计动作价值函数 Q ，它提供了智能体从任意给定状态遵循其策略所能获得的折扣奖励的估计值。我们在训练开始前通过运行随机策略收集一组固定状态，并跟踪这些状态的最大预测值²的平均值 Q 。图2中最右侧的两个图表显示，平均预测值 Q 的增长比智能体获得的平均总奖励平滑得多，且在其他五款游戏中绘制相同指标也产生了类似的平滑曲线。除了观察到训练期间预测值 Q 相对平滑的改进外，我们在所有实验中均未遇到任何发散问题。这表明，尽管缺乏理论收敛保证，我们的方法能够以稳定的方式，利用强化学习信号和随机梯度下降训练大型神经网络。

²The maximum for each state is taken over the possible actions.




Figure 2: The two plots on the left show average reward per episode on Breakout and Seaquest respectively during training. The statistics were computed by running an ϵ -greedy policy with $\epsilon = 0.05$ for 10000 steps. The two plots on the right show the average maximum predicted action-value of a held out set of states on Breakout and Seaquest respectively. One epoch corresponds to 50000 minibatch weight updates or roughly 30 minutes of training time.



Figure 3: The leftmost plot shows the predicted value function for a 30 frame segment of the game Seaquest. The three screenshots correspond to the frames labeled by A, B, and C respectively.

5.2 Visualizing the Value Function

Figure 3 shows a visualization of the learned value function on the game Seaquest. The figure shows that the predicted value jumps after an enemy appears on the left of the screen (point A). The agent then fires a torpedo at the enemy and the predicted value peaks as the torpedo is about to hit the enemy (point B). Finally, the value falls to roughly its original value after the enemy disappears (point C). Figure 3 demonstrates that our method is able to learn how the value function evolves for a reasonably complex sequence of events.

5.3 Main Evaluation

We compare our results with the best performing methods from the RL literature [3, 4]. The method labeled **Sarsa** used the Sarsa algorithm to learn linear policies on several different feature sets hand-engineered for the Atari task and we report the score for the best performing feature set [3]. **Contingency** used the same basic approach as **Sarsa** but augmented the feature sets with a learned representation of the parts of the screen that are under the agent’s control [4]. Note that both of these methods incorporate significant prior knowledge about the visual problem by using background subtraction and treating each of the 128 colors as a separate channel. Since many of the Atari games use one distinct color for each type of object, treating each color as a separate channel can be similar to producing a separate binary map encoding the presence of each object type. In contrast, our agents only receive the raw RGB screenshots as input and must *learn* to detect objects on their own.

In addition to the learned agents, we also report scores for an expert human game player and a policy that selects actions uniformly at random. The human performance is the median reward achieved after around two hours of playing each game. Note that our reported human scores are much higher than the ones in Bellemare et al. [3]. For the learned methods, we follow the evaluation strategy used in Bellemare et al. [3, 5] and report the average score obtained by running an ϵ -greedy policy with $\epsilon = 0.05$ for a fixed number of steps. The first five rows of table 1 show the per-game average scores on all games. Our approach (labeled DQN) outperforms the other learning methods by a substantial margin on all seven games despite incorporating almost no prior knowledge about the inputs.

We also include a comparison to the evolutionary policy search approach from [8] in the last three rows of table 1. We report two sets of results for this method. The **HNeat Best** score reflects the results obtained by using a hand-engineered object detector algorithm that outputs the locations and




图2: 左侧两图分别展示了训练期间Breakout和Seaquest游戏每回合的平均奖励。这些统计数据是通过运行一个 ϵ -贪婪策略计算得出的, 其中 ϵ = 0.05, 共执行了10000步。右侧两图分别展示了Breakout和Seaquest游戏中一组保留状态的平均最大预测动作值。一个训练周期对应50000次小批量权重更新, 约等于30分钟的训练时间。



图3: 最左侧的图表展示了游戏《海底战争》一段30帧片段的价值函数预测值。三张屏幕截图分别对应标注为A、B和C的帧。

5.2 价值函数可视化

图3展示了在游戏《海底战争》中学到的价值函数的可视化。该图显示, 当敌人在屏幕左侧出现时(点A), 预测价值会跃升。随后, 代理向敌人发射鱼雷, 当鱼雷即将击中敌人时(点B), 预测价值达到峰值。最后, 在敌人消失后(点C), 价值大致回落至原始水平。图3表明, 我们的方法能够学习价值函数如何在一个相对复杂的事件序列中演变。

5.3 主要评估

我们将自己的结果与强化学习文献中表现最佳的方法进行了比较[3, 4]。标记为Sarsa的方法使用Sarsa算法学习针对Atari任务手工设计的多个不同特征集上的线性策略, 我们报告了表现最佳特征集的得分[3]。Contingency采用与Sarsa相同的基础方法, 但通过添加对智能体控制下屏幕部分的学习表示来增强特征集[4]。需要注意的是, 这两种方法都通过使用背景差分并将128种颜色中的每一种视为独立通道, 融入了关于视觉问题的显著先验知识。由于许多Atari游戏为每种对象类型使用一种独特颜色, 将每种颜色视为独立通道类似于生成编码每种对象类型存在的独立二值图。相比之下, 我们的智能体仅接收原始RGB屏幕截图作为输入, 必须*learn*自行检测对象。

除了学习型智能体外, 我们还报告了一位专业人类游戏玩家和随机均匀选择动作策略的得分。人类表现是每款游戏游玩约两小时后获得的中位数奖励。需注意, 我们报告的人类得分远高于Bellemare等人[3]的研究结果。对于学习方法, 我们遵循Bellemare等人[3,5]采用的评估策略, 报告通过运行 ϵ -贪婪策略(ϵ = 0.05)在固定步数内获得的平均得分。表1前五行展示了所有游戏的单局平均得分。尽管几乎未引入任何输入先验知识, 我们的方法(标记为DQN)在所有七款游戏上均以显著优势超越其他学习方法。

我们还在表1的最后三行中加入了与[8]中进化策略搜索方法的比较。我们为此方法报告了两组结果。HNeat最佳分数反映了使用手动设计的物体检测算法所获得的结果, 该算法输出位置

	B. Rider	Breakout	Enduro	Pong	Q*bert	Seaquest	S. Invaders
Random	354	1.2	0	−20.4	157	110	179
Sarsa [3]	996	5.2	129	−19	614	665	271
Contingency [4]	1743	6	159	−17	960	723	268
DQN	4092	168	470	20	1952	1705	581
Human	7456	31	368	−3	18900	28010	3690
HNeat Best [8]	3616	52	106	19	1800	920	1720
HNeat Pixel [8]	1332	4	91	−16	1325	800	1145
DQN Best	5184	225	661	21	4500	1740	1075

Table 1: The upper table compares average total reward for various learning methods by running an ϵ -greedy policy with $\epsilon = 0.05$ for a fixed number of steps. The lower table reports results of the single best performing episode for HNeat and DQN. HNeat produces deterministic policies that always get the same score while DQN used an ϵ -greedy policy with $\epsilon = 0.05$.

types of objects on the Atari screen. The **HNeat Pixel** score is obtained by using the special 8 color channel representation of the Atari emulator that represents an object label map at each channel. This method relies heavily on finding a deterministic sequence of states that represents a successful exploit. It is unlikely that strategies learnt in this way will generalize to random perturbations; therefore the algorithm was only evaluated on the highest scoring single episode. In contrast, our algorithm is evaluated on ϵ -greedy control sequences, and must therefore generalize across a wide variety of possible situations. Nevertheless, we show that on all the games, except Space Invaders, not only our max evaluation results (row 8), but also our average results (row 4) achieve better performance.

Finally, we show that our method achieves better performance than an expert human player on Breakout, Enduro and Pong and it achieves close to human performance on Beam Rider. The games Q*bert, Seaquest, Space Invaders, on which we are far from human performance, are more challenging because they require the network to find a strategy that extends over long time scales.

6 Conclusion

This paper introduced a new deep learning model for reinforcement learning, and demonstrated its ability to master difficult control policies for Atari 2600 computer games, using only raw pixels as input. We also presented a variant of online Q-learning that combines stochastic minibatch updates with experience replay memory to ease the training of deep networks for RL. Our approach gave state-of-the-art results in six of the seven games it was tested on, with no adjustment of the architecture or hyperparameters.

References

- [1] Leemon Baird. Residual algorithms: Reinforcement learning with function approximation. In *Proceedings of the 12th International Conference on Machine Learning (ICML 1995)*, pages 30–37. Morgan Kaufmann, 1995.
- [2] Marc Bellemare, Joel Veness, and Michael Bowling. Sketch-based linear value function approximation. In *Advances in Neural Information Processing Systems 25*, pages 2222–2230, 2012.
- [3] Marc G Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. The arcade learning environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research*, 47:253–279, 2013.
- [4] Marc G Bellemare, Joel Veness, and Michael Bowling. Investigating contingency awareness using atari 2600 games. In *AAAI*, 2012.
- [5] Marc G. Bellemare, Joel Veness, and Michael Bowling. Bayesian learning of recursively factored environments. In *Proceedings of the Thirtieth International Conference on Machine Learning (ICML 2013)*, pages 1211–1219, 2013.

	B. Rider	Breakout	Enduro	Pong	Q*bert	Seaquest	S. Invaders
Random	354	1.2	0	−20.4	157	110	179
Sarsa [3]	996	5.2	129	−19	614	665	271
Contingency [4]	1743	6	159	−17	960	723	268
DQN	4092	168	470	20	1952	1705	581
Human	7456	31	368	−3	18900	28010	3690
HNeat Best [8]	3616	52	106	19	1800	920	1720
HNeat Pixel [8]	1332	4	91	−16	1325	800	1145
DQN Best	5184	225	661	21	4500	1740	1075

表1: 上表通过运行一个 ϵ -贪婪策略，其中 $\epsilon = 0.05$ ，在固定步数内比较了各种学习方法的平均总奖励。下表报告了HNeat和DQN在单次最佳表现回合中的结果。HNeat产生确定性策略，总是获得相同的分数，而DQN使用了 ϵ -贪婪策略，其中 $\epsilon = 0.05$ 。

雅达利屏幕上的物体类型。HNeat像素分数是通过使用雅达利模拟器的特殊8色通道表示获得的，该表示在每个通道上呈现物体标签图。这种方法严重依赖于寻找代表成功利用的确定性状态序列。以这种方式学习的策略不太可能推广到随机扰动；因此，该算法仅在得分最高的单次游戏中进行评估。相比之下，我们的算法在 ϵ -贪婪控制序列上进行评估，因此必须推广到各种可能的情况。尽管如此，我们表明，在除《太空侵略者》外的所有游戏中，不仅我们的最大评估结果（第8行），而且我们的平均结果（第4行）都实现了更好的性能。

最后，我们展示了我们的方法在《打砖块》、《拉力赛》和《乒乓球》上取得了比人类专家玩家更好的表现，并且在《光束骑士》上接近人类水平。而在《Q*bert》、《海底冒险》和《太空侵略者》这些游戏中，我们的表现还远不及人类，这些游戏更具挑战性，因为它们要求网络找到一种能在长时间尺度上持续有效的策略。

6 结论

本文介绍了一种用于强化学习的新型深度学习模型，并展示了其仅使用原始像素作为输入即可掌握Atari 2600电脑游戏中复杂控制策略的能力。我们还提出了一种在线Q学习的变体，它将随机小批量更新与经验回放记忆相结合，以简化强化学习中深度网络的训练。我们的方法在测试的七款游戏中的六款上取得了最先进的结果，且无需调整架构或超参数。

参考文献

- [1] Leemon Baird. 残差算法：基于函数逼近的强化学习。载于 *Proceedings of the 12th International Conference on Machine Learning (ICML 1995)*，第30–37页。Morgan Kaufmann, 1995年。[2] Marc Bellemare、Joel Veness 和 Michael Bowling。基于草图的线性价值函数逼近。载于 *Advances in Neural Information Processing Systems 25*，第2222–2230页，2012年。[3] Marc G Bellemare、Yavar Naddaf、Joel Veness 和 Michael Bowling。街机学习环境：通用智能体的评估平台。 *Journal of Artificial Intelligence Research*，第47卷，第253–279页，2013年。[4] Marc G Bellemare、Joel Veness 和 Michael Bowling。使用 Atari 2600 游戏研究应急感知。载于 *AAAI*，2012年。[5] Marc G. Bellemare、Joel Veness 和 Michael Bowling。递归分解环境的贝叶斯学习。载于 *Proceedings of the Thirtieth International Conference on Machine Learning (ICML 2013)*，第1211–1219页，2013年。

- [6] George E. Dahl, Dong Yu, Li Deng, and Alex Acero. Context-dependent pre-trained deep neural networks for large-vocabulary speech recognition. *Audio, Speech, and Language Processing, IEEE Transactions on*, 20(1):30–42, January 2012.
- [7] Alex Graves, Abdel-rahman Mohamed, and Geoffrey E. Hinton. Speech recognition with deep recurrent neural networks. In *Proc. ICASSP*, 2013.
- [8] Matthew Hausknecht, Risto Miikkulainen, and Peter Stone. A neuro-evolution approach to general atari game playing. 2013.
- [9] Nicolas Heess, David Silver, and Yee Whye Teh. Actor-critic reinforcement learning with energy-based policies. In *European Workshop on Reinforcement Learning*, page 43, 2012.
- [10] Kevin Jarrett, Koray Kavukcuoglu, MarcAurelio Ranzato, and Yann LeCun. What is the best multi-stage architecture for object recognition? In *Proc. International Conference on Computer Vision and Pattern Recognition (CVPR 2009)*, pages 2146–2153. IEEE, 2009.
- [11] Alex Krizhevsky, Ilya Sutskever, and Geoff Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems 25*, pages 1106–1114, 2012.
- [12] Sascha Lange and Martin Riedmiller. Deep auto-encoder neural networks in reinforcement learning. In *Neural Networks (IJCNN), The 2010 International Joint Conference on*, pages 1–8. IEEE, 2010.
- [13] Long-Ji Lin. Reinforcement learning for robots using neural networks. Technical report, DTIC Document, 1993.
- [14] Hamid Maei, Csaba Szepesvari, Shalabh Bhatnagar, Doina Precup, David Silver, and Rich Sutton. Convergent Temporal-Difference Learning with Arbitrary Smooth Function Approximation. In *Advances in Neural Information Processing Systems 22*, pages 1204–1212, 2009.
- [15] Hamid Maei, Csaba Szepesvári, Shalabh Bhatnagar, and Richard S. Sutton. Toward off-policy learning control with function approximation. In *Proceedings of the 27th International Conference on Machine Learning (ICML 2010)*, pages 719–726, 2010.
- [16] Volodymyr Mnih. *Machine Learning for Aerial Image Labeling*. PhD thesis, University of Toronto, 2013.
- [17] Andrew Moore and Chris Atkeson. Prioritized sweeping: Reinforcement learning with less data and less real time. *Machine Learning*, 13:103–130, 1993.
- [18] Vinod Nair and Geoffrey E Hinton. Rectified linear units improve restricted boltzmann machines. In *Proceedings of the 27th International Conference on Machine Learning (ICML 2010)*, pages 807–814, 2010.
- [19] Jordan B. Pollack and Alan D. Blair. Why did td-gammon work. In *Advances in Neural Information Processing Systems 9*, pages 10–16, 1996.
- [20] Martin Riedmiller. Neural fitted q iteration—first experiences with a data efficient neural reinforcement learning method. In *Machine Learning: ECML 2005*, pages 317–328. Springer, 2005.
- [21] Brian Sallans and Geoffrey E. Hinton. Reinforcement learning with factored states and actions. *Journal of Machine Learning Research*, 5:1063–1088, 2004.
- [22] Pierre Sermanet, Koray Kavukcuoglu, Soumith Chintala, and Yann LeCun. Pedestrian detection with unsupervised multi-stage feature learning. In *Proc. International Conference on Computer Vision and Pattern Recognition (CVPR 2013)*. IEEE, 2013.
- [23] Richard Sutton and Andrew Barto. *Reinforcement Learning: An Introduction*. MIT Press, 1998.
- [24] Gerald Tesauero. Temporal difference learning and td-gammon. *Communications of the ACM*, 38(3):58–68, 1995.
- [25] John N Tsitsiklis and Benjamin Van Roy. An analysis of temporal-difference learning with function approximation. *Automatic Control, IEEE Transactions on*, 42(5):674–690, 1997.
- [26] Christopher JCH Watkins and Peter Dayan. Q-learning. *Machine learning*, 8(3-4):279–292, 1992.

- [6] George E. Dahl, Dong Yu, Li Deng, and Alex Acero. 面向大词汇量语音识别的上下文相关预训练深度神经网络。 *Audio, Speech, and Language Processing, IEEE Transactions on*, 20(1):30–42, 2012年1月。
- [7] Alex Graves, Abdel-rahman Mohamed, and Geoffrey E. Hinton. 基于深度循环神经网络的语音识别。收录于 *Proc. ICASSP*, 2013年。
- [8] Matthew Hausknecht, Risto Miikkulainen, and Peter Stone. 一种面向通用Atari游戏对战的神经进化方法。2013年。
- [9] Nicolas Heess, David Silver, and Yee Whye Teh. 基于能量策略的行动者-评论家强化学习。收录于 *European Workshop on Reinforcement Learning*, 第43页, 2012年。
- [10] Kevin Jarrett, Koray Kavukcuoglu, MarcAurelio Ranzato, and Yann LeCun. 什么是目标识别的最佳多阶段架构? 收录于 *Proc. International Conference on Computer Vision and Pattern Recognition (CVPR 2009)*, 第2146–2153页。IEEE, 2009年。
- [11] Alex Krizhevsky, Ilya Sutskever, and Geoff Hinton. 基于深度卷积神经网络的ImageNet分类。收录于 *Advances in Neural Information Processing Systems 25*, 第1106–1114页, 2012年。
- [12] Sascha Lange and Martin Riedmiller. 强化学习中的深度自编码器神经网络。收录于 *Neural Networks (IJCNN), The 2010 International Joint Conference on*, 第1–8页。IEEE, 2010年。
- [13] Long-Ji Lin. 机器人使用神经网络的强化学习。技术报告, DTIC Document, 1993年。
- [14] Hamid Maei, Csaba Szepesvari, Shalabh Bhatnagar, Doina Precup, David Silver, and Rich Sutton. 任意平滑函数逼近下的收敛时序差分学习。收录于 *Advances in Neural Information Processing Systems 22*, 第1204–1212页, 2009年。
- [15] Hamid Maei, Csaba Szepesvári, Shalabh Bhatnagar, and Richard S. Sutton. 面向函数逼近的离策略学习控制。收录于 *Proceedings of the 27th International Conference on Machine Learning (ICML 2010)*, 第719–726页, 2010年。
- [16] Volodymyr Mnih. *Machine Learning for Aerial Image Labeling*. 博士论文, 多伦多大学, 2013年。
- [17] Andrew Moore and Chris Atkeson. 优先扫描: 数据更少、实时性更低的强化学习。 *Machine Learning*, 13:103–130, 1993年。
- [18] Vinod Nair and Geoffrey E Hinton. 修正线性单元改进受限玻尔兹曼机。收录于 *Proceedings of the 27th International Conference on Machine Learning (ICML 2010)*, 第807–814页, 2010年。
- [19] Jordan B. Pollack and Alan D. Blair. TD-Gammon为何有效。收录于 *Advances in Neural Information Processing Systems 9*, 第10–16页, 1996年。
- [20] Martin Riedmiller. 神经拟合Q迭代——一种数据高效神经强化学习方法的初步经验。收录于 *Machine Learning: ECML 2005*, 第317–328页。Springer, 2005年。
- [21] Brian Sallans and Geoffrey E. Hinton. 基于因子化状态与动作的强化学习。 *Journal of Machine Learning Research*, 5:1063–1088, 2004年。
- [22] Pierre Sermanet, Koray Kavukcuoglu, Soumith Chintala, and Yann LeCun. 基于无监督多阶段特征学习的行人检测。收录于 *Proc. International Conference on Computer Vision and Pattern Recognition (CVPR 2013)*。IEEE, 2013年。
- [23] Richard Sutton and Andrew Barto. *Reinforcement Learning: An Introduction*. MIT Press, 1998年。
- [24] Gerald Tesauero. 时序差分学习与TD-Gammon。 *Communications of the ACM*, 38(3):58–68, 1995年。
- [25] John N Tsitsiklis and Benjamin Van Roy. 函数逼近下时序差分学习的分析。 *Automatic Control, IEEE Transactions on*, 42(5):674–690, 1997年。
- [26] Christopher JCH Watkins and Peter Dayan. Q学习。 *Machine learning*, 8(3-4):279–292, 1992年。